

# Domain 1

## Network Connectivity Strategies

### AWS Direct Connect vs. AWS VPN vs. Transit Gateway

#### 1. AWS Direct Connect

- **Purpose:** Provides a dedicated, private connection between on-premises data centers and AWS.
- **Performance:** Offers stable, low-latency, and high-bandwidth connections, allowing consistent throughput.
- **Security:** Private connectivity ensures data doesn't traverse the public Internet, enhancing security.
- **Use Cases:** Ideal for data-intensive applications, hybrid cloud setups, and regulatory compliance.

#### 2. AWS VPN

- **Purpose:** Establishes a secure, encrypted connection over the public Internet between on-premises networks and AWS VPCs.
- **Performance:** Latency can vary as it relies on Internet traffic; generally lower bandwidth than Direct Connect.
- **Security:** Utilizes IPsec for encryption, ensuring secure communication over untrusted networks.
- **Use Cases:** Suitable for smaller data transfers, temporary connections, and when budgets are tight.

#### 3. AWS Transit Gateway

- **Purpose:** Acts as a central hub to connect multiple VPCs and on-premises networks.
- **Performance:** Optimizes routing and reduces the need for multiple connections, potentially enhancing network performance.
- **Scalability:** Simplifies management by allowing multiple VPCs to connect through a single gateway; ideal for large, multi-region architectures.
- **Use Cases:** Best for complex architectures with many VPCs, requiring centralized traffic management.

### Summary

- Use **Direct Connect** for high-performance, stable connections.
  - Choose **VPN** for secure, encrypted Internet connections, especially for smaller workloads.
  - Implement **Transit Gateway** for efficient management of multiple VPCs and connectivity in expansive architectures.
-

# Domain 2

## Compute & Serverless Architectures

### EC2 Auto Scaling vs. Lambda Provisioning

#### 1. Scalability

- **EC2 Auto Scaling:**

- Scales EC2 instances up or down based on predefined metrics (CPU, memory, etc.).
- Requires time to spin up/down instances, making it less responsive to sudden spikes.

- **Lambda:**

- Automatically scales instantly based on incoming requests.
- Handles unpredictable traffic seamlessly, as it can manage thousands of concurrent executions.

#### 2. Cost Management

- **EC2 Auto Scaling:**

- Pay for running instances, regardless of usage.
- Costs can accumulate during idle periods, depending on instance types.

- **Lambda:**

- Pay only for execution time and requests.
- More cost-effective for sporadic workloads due to a consumption-based pricing model.

#### 3. Management

- **EC2 Auto Scaling:**

- Requires more management overhead (instance health checks, AMI maintenance).
- Infrastructure management such as OS updates, scaling policies, etc.

- **Lambda:**

- Fully managed service; no infrastructure management required.
- Simplifies deployments, updates, and maintenance.

#### 4. Performance

- **EC2 Auto Scaling:**

- Better suited for long-running applications that need consistent performance.
- Initialization overhead can lead to latency during scale-up events.

- **Lambda:**

- Optimal for short-lived tasks; may experience cold starts, impacting latency.
- Excellent for event-driven architectures with variable workloads.

## 5. Use Cases

- **EC2 Auto Scaling:**

- Best for predictable scaling of applications like web servers, databases, and services with steady workloads.

- **Lambda:**

- Ideal for microservices, APIs, real-time data processing, and event-driven applications.

## Conclusion

For unpredictable traffic patterns, **Lambda** often provides better scalability and cost efficiency, while **EC2 Auto Scaling** may be preferable for stable, long-running workloads. The choice depends on specific application needs and usage patterns.

---